

Introduction to Linux and Supercomputers

Zhores Software (modules, batch)

MA030366 (Term 2) R2-B5-2026

Zacharov Igor

i.zacharov@skoltech.ru

SLURM

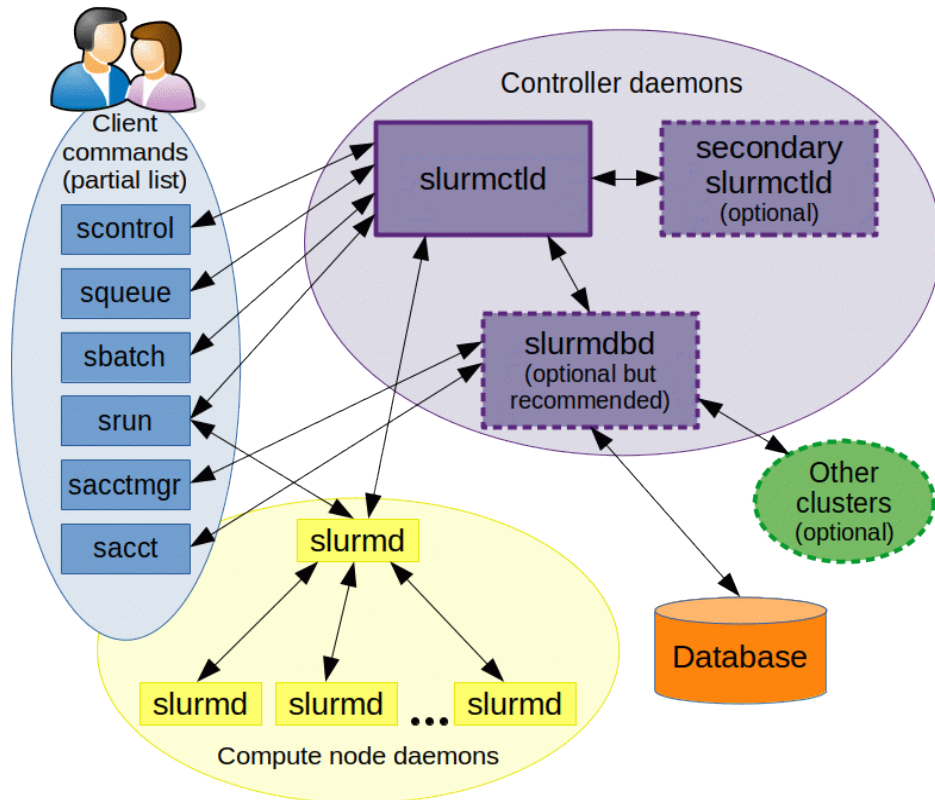
Simple Linux Utility for Resource Management

- Job scheduling in cluster systems
 - Job: self-contained unit of work
 - Work may be described (automated) in a (bash) script
 - Scheduling: an upfront allocation of resources necessary to complete the job.
 - Scheduling may be based on:
 - Job priority
 - Estimated execution time
 - Resource requirements (number of CPU/GPU needed, size memory, etc.)
 - Job Queue: data structure maintained by Scheduler containing jobs to run
 - Submit job: action of copying job (description = script) in to the queue (data structure)
- Use of batch queue gives these benefits:
 - Sharing of computer resources among many users
 - Time-shift processing when resources are available
 - Avoids human supervision for scheduling
 - Allows around-the-clock utilization of computing resources
- Other such systems
 - NQE/PBS, Condor, LSF, SUN Grid Engine (SGE)

SLURM Queues Architecture

- Partition
 - Group of nodes with specific characteristic
 - Queue is organizing access of jobs for a partition
 - Job may have multiple steps, which are sets of tasks within a job
- Slurm does four basic steps to manage CPU resources for job/step:
 1. Selection of nodes
 2. Allocation of CPUs from selected nodes
 3. Distribution of tasks to selected nodes
 4. Optional binding of tasks to allocated CPUs

- One central controller daemon (slurmctld) on management node
 - Obtains resource utilization information
 - Makes scheduling decisions within partitions
- A daemon on each computing node (slurmd)
- One central daemon for the accounting database (slurmdbd)



SLURM Job status change

- SLURM provides the Final-State machine for changing the job status
 - At submission it is “PENDING” (PD) – waiting for the resources
 - When resources are allocated the job is “RUNNING” (R)
 - The status is given in the output of the **squeue [-l]** command

Status	Code	Explanation
COMPLETED	CD	Job has completed
COMPLETING	CG	Job is finishing but some processes are still active
FAILED	F	Job terminated with a non-zero exit code (failed)
PENDING	PD	Job is waiting for resource allocation
PREEMPTED	PR	Job was terminated because of preemption by another job
RUNNING	R	Job currently is allocated to a node and running
SUSPENDED	S	Job has been stopped with cores released to another job
STOPPED	ST	Job has been stopped with its cores retained

JOBID	PARTITION	NAME	USER	STATE	TIME	TIME_LIMI	NODES	NODELIST(REASON)
1169972	gpu_devel	bash	i.zachar	COMPLETING	1:13	1:00	1	vt01

Zhores Queue (partition) structure

- The queue (partition) composition on Zhores:
 - Special partitions:
 - QUEUES: res (gn[21-25]), chess (cn[01-16], gn[16-20], ct[01-04]), gpu_a100 (gn26)
 - Available to all users:

Characteristic/Queue	cpu	gpu	mem	gpu_devel	htc
# nodes per task	unlimited	unlimited	unlimited	2	1
Max processing time /default time	6/1 day	6/1day	6/1 day	12/- h	1/- day
Oversubscribe	no	no	no	2	no
QoS	cpu	gpu	mem	gpu_devel	htc
#nodes/CPU's	32/768	15/540	6/480	14/56	83/2852
Nodes	cn[17-44] hd[01-04]	gn[01-15]	ct[05-10]	vt[01-14]	ct[01-10] cn[01-44] hd[01-04] gn[01-25]

No sharing
of resources

- Need to specify the parameters to fully allocate job resources in a queue:
 - To submit into the cpu partition a job expected to run 45 minutes on 4 nodes using 16 cores each and 100 GB memory used per node:

```
sbatch --partition=cpu --nodes=4 -ntasks=16 --mem=100G --time=45:00 runscript.sh
```

- Same can be done within the script:

- Ex.: sbatch runscript.sh

```
#SBATCH -p cpu
#SBATCH -N 1
#SBATCH -n 16
#SBATCH --mem=100G
#SBATCH -t 45:00
```

SLURM - LAB

- Admin commands to see configuration, queues, etc:
 - Composition of queues
 - sinfo
 - scontrol show partitions
 - squeue (as root to see all jobs)
- To allocate resources for running jobs (semi-)interactively

– `salloc -p gpu_devel -t 5 -N 1 -n 1`

– `srun [-w vt11] x.sh`

– `srun ls`

– `srun hostname`

– `squeue [-j 1169093]`

– `scontrol show job 1169093`

– `scontrol show nodes vt11`

```
#!/bin/bash
hostname
date
sleep 45
```

```
bash-4.2$ salloc -p gpu_devel -t 5 -N 1 -n 1
salloc: Pending job allocation 1169093
salloc: job 1169093 queued and waiting for resources
salloc: job 1169093 has been allocated resources
salloc: Granted job allocation 1169093
salloc: Waiting for resource configuration
salloc: Nodes vt11 are ready for job
```

```
bash-4.2$ sbatch -p gpu_devel -t 1 -N 1 -n 1 x.sh
```

```
Submitted batch job 1169095
```

```
bash-4.2$ squeue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
1169095	gpu_devel	x.sh	i.zachar	R	0:02	1	vt11

```
Bash-4.2$ ls
```

```
-rwxrwxr-x 1 i.zacharov i.zacharov 44 Oct 24 20:00 x.sh
```

```
-rw-rw-r-- 1 i.zacharov i.zacharov 41 Oct 24 20:01 slurm-1169095.out
```


Output file

Uses previously
allocated resources

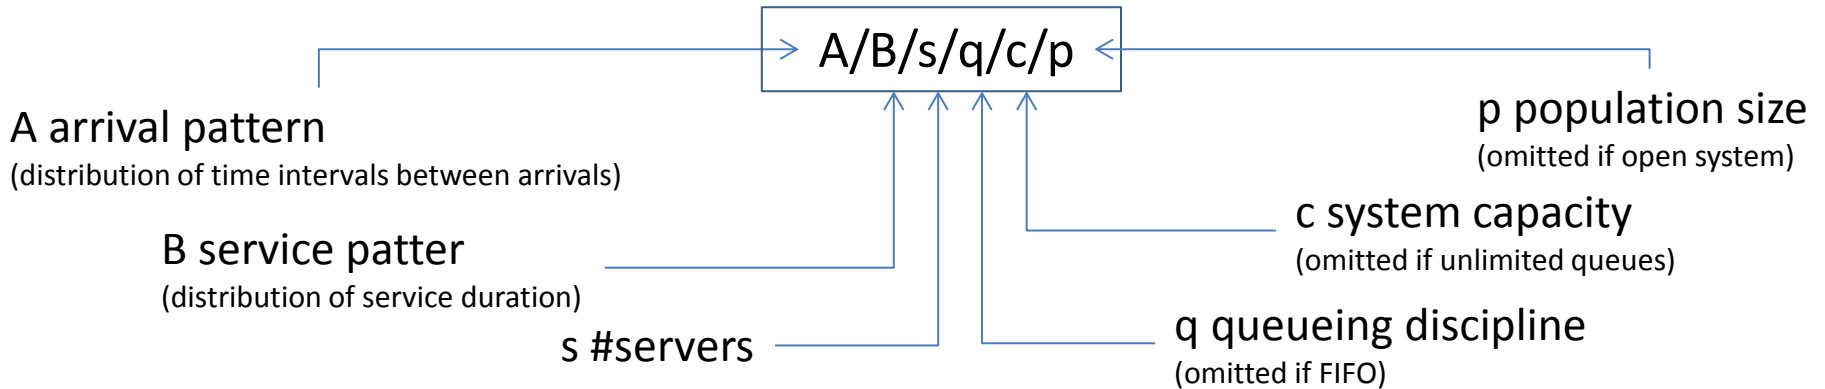
tasks

Theory: Elements of a Queueing System

- Users compete for the resources (services) by generating jobs (requests)
 - Potentially unlimited #jobs (open system), but in practice some bound number (closed system)
- Arrival of jobs in the queue at time t
 - Interval between adjacent arrivals is randomly distributed
- Queue represents jobs waiting for service
 - It can be empty
 - Job in service is not in the queue
 - Queue maximum size: maximum number jobs waiting = System capacity
 - Queueing discipline: organization of inserting to/from queue
 - FIFO = First In First Out (also FCFS = First Come First Serve) – orderly queue
 - LIFO = Last In First Out (also LCFS = Last Come First Serve) - stack
 - SIRO = Serve In Random Order
 - Priority Queue – typically number of queues with different priorities
 - Complex scheduling discipline where jobs change their position based on time spent in q , availability of the service (i.e. “backfill”), etc.
- Service represents activity that takes (processing) time
 - For each job it may take different time (random)
- Queueing Theory gives system performance as function of input parameters
 - Average waiting time as function of the average service time & #arrivals
 - Throughput (#jobs completed per time) as function of average service time & #arrivals



Kendall Classification of Queueing Systems



- A and B can be a
 - Poisson (Markovian time) distribution : M
 - Erlang distribution: E
 - Deterministic arrivals and/or constant service duration
 - General (any) distribution
- Zhores queueing system is $M/M/n$, where n is the number of servers in a queue
 - Zhores has several queues, the analysis if for each queue separately
 - Arrival rate λ [$\frac{1}{s}$], service rate μ [$\frac{1}{s}$], **utilization** $\rho = \frac{\lambda}{n\mu} < 1$ (if >1 the queue will grow no bound)
 - Probability for job to join the queue (all n servers occupied)
 - Average #jobs in the system:
 - Average response time (waiting time in the system):

$$N_s = \frac{\rho}{1 - \rho} C(n, \lambda/\mu) + n\rho$$

$$C(n, \lambda/\mu) = \frac{1}{1 + (1 - \rho) \left(\frac{n!}{(n\rho)^n} \right) \sum_{k=0}^{n-1} \frac{(n\rho)^k}{k!}}$$

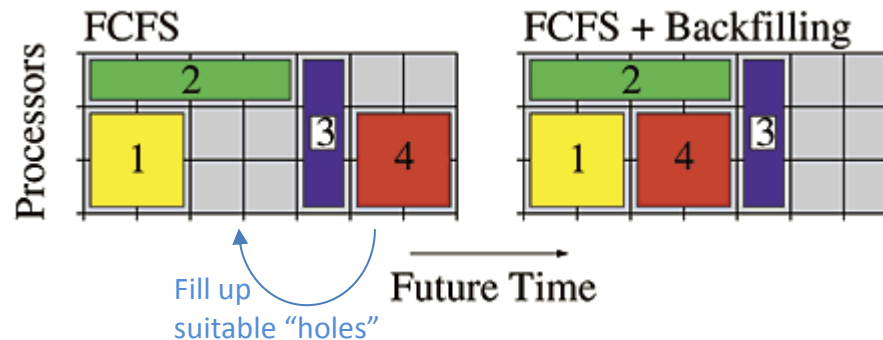
$$t_w = \frac{C(n, \lambda/\mu)}{n\mu - \lambda} + \frac{1}{\mu}$$

Queueing Theory: conclusions for Zhores

- Estimate resources for your job for best turnaround
 - Specify in sbatch command or job script:
 - Partition, Runtime, #nodes, memory needed, optimal CPU/GPU or task number
- Fastest queue is the queue with smaller average processing time
 - Waiting time is $\sim 1/\mu$
 - -> Prefer the htc queue for faster turnaround
 - Expect that there could be many jobs in the queue (~ 100 s of jobs)
- Specify expected Runtime accurately
 - The scheduler *Backfill strategy* may allow you to get results faster
 - Schedule jobs as long as they don't delay a waiting job that is higher in the queue
 - Increases utilization of the cluster (admins will like you)

$$t_w \sim \frac{1}{n\mu - \lambda} + \frac{1}{\mu}$$

$$N_s \sim \frac{\rho}{1 - \rho} + n\rho, \quad \text{while } \rho \approx 1$$



Slurm Interactive jobs - Lab

- Access to an allocated node interactively
 - Eg. for testing specialized (parallel) software that cannot be done on the access node
 - Performance testing: the allocated resources will not be shared

```
srun --time=1:00:00 -p gpu_devel -N 1 -n 1 --pty bash -l
```

- srun: job 1169972 queued and waiting for resources
- srun: job 1169972 has been allocated resources

- Checking: `squeue -l`

JOBID	PARTITION	NAME	USER	STATE	TIME	TIME_LIMI	NODES	NODELIST(REASON)
1169972	gpu_devel	bash	i.zachar	COMPLETING	1:13	1:00	1	vt01

- Difference between `salloc` + `srun` and standalone `srun`:

- The `salloc/sbatch` allocate resources

- The `srun` inherits these resources and executes commands as job steps
- Multiple `srun` within a single allocation is possible (multiple job steps)
- Check job steps with: `salloc -p htc -N 1 -n 1 -c 4 -t 5 -mem=1G`
`./piloop.sh`
`sacct -j $SLURM_JOB_ID`
- (this works on Zhores and doesn't work in the sandbox)
- **This is sometimes called a "Packed Job"**

```
#piloop.sh:  
# salloc -N 1 -n 1  
for ((i=1; i<8; i++))  
do  
    srun runPlomp 10 &  
done
```

- The `srun` standalone allocates resources and performs one job step
- Both, `salloc` and `srun` block until resources are available (interactive)
- The `sbatch` queues the job waiting for the resource allocation

SLURM parallel jobs

- Models for Parallel job = tasks running simultaneously
 - A multi-process program
 - Single Process, Multiple Data (SPMD) paradigm, eg. with MPI
 - A multi-threaded program
 - Shared memory paradigm, eg. with OpenMP or pthreads
 - Several instances of a single-threaded program
 - Embarrassingly parallel paradigm, implemented as a job array
 - One master program controlling several slave programs
 - Master/Slave paradigm
- SLURM
 - A task represents a process
 - Request with `–ntasks`
 - Can be split on several compute nodes
 - A multi-process program is made of several tasks
 - A multi-threaded program is composed on only one task
 - Uses several CPUs, request with `–cpus-per-task` (or `–c` in short)
 - Cannot be split across several compute nodes

SLURM OpenMP example

- OpenMP program: make piomp
- Bash script: runomp.sh
 - Start with: sbatch runomp.sh
 - Or: sbatch -c X runomp.sh
 - Most important parameter: -c setting number of cpus per task
 - Note: -N 1 (--nodes)

```
#!/usr/bin/bash
#SBATCH --job-name=piomp
#SBATCH --output=piomp_%j.out
#SBATCH --partition=cpu          # -p cpu
#SBATCH --nodes=1                # -N 1 number of nodes
#SBATCH --cpus-per-task=8        # -c X cpus assigned to each task
#SBATCH --time=2:00
#SBATCH --mem=1G                 # memory per node

export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK:-1}

module load mpi/openmpi-3.1.2
module load compilers/gcc-7.3.1
make piomp &>/dev/null

mtrials=${1:-1}
numtrials=$mtrials
seed=$RANDOM

/usr/bin/time -p ./piomp -t $numtrials -a -t $numtrials -s $seed
```

SLURM MPI example

- MPI program: make pimpi
- Bash script: runmpi.sh
 - Start with: sbatch runmpi.sh
 - Note: -N (--nodes) = -n (ntasks) for distribution of MPI processes to nodes
 - Optimization: leave N unspecified (lump MPI processes on the same nodes)
 - Try -n x -N x/2
 - mpirun “-np auto”
 - from resources

```
#!/usr/bin/bash
#SBATCH --job-name=pimpi
#SBATCH --output=pimpi_%j.out
#SBATCH --partition=cpu          # -p cpu
#SBATCH --nodes=4                # -N 4 number of nodes
#SBATCH --ntasks=4               # -n 4 total number of mpi processes
#SBATCH --cpus-per-task=8        # -c X cpus assigned to each task
#SBATCH --mincpus=9              # to ensure sufficient resources for MPI master
#SBATCH --time=2:00
#SBATCH --mem=1G                 # memory per node

export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK:-1}
export MPI_NUM_PROCESS=${SLURM_JOB_NUM_NODES:-4}

module load mpi/openmpi-3.1.2
module load compilers/gcc-7.3.1
make pimpi &>/dev/null

mtrials=${1:-1}
numtrials=$mtrials
seed=$RANDOM

/usr/bin/time -p mpirun ./pimpi -a -t $numtrials -s $seed
```

SLURM Embarrassingly parallel example

- Multiple programs running same code with different parameter(s)
 - Same OpenMP program: make piomp will run with OMP_NUM_THREADS=1
 - Can also be compiled without -fopenmp, only scalar code is needed
- Bash script: runarr.sh
- SBATCH -a 1-N%X
 - X is #running same time
- Output is collected in log file (eg. “tee -a log”)
- Analyze the log to assemble the result
 - Ex.: analyze.sh
 - Awk is used in this case
 - Python could be used instead

```
#!/usr/bin/bash
#SBATCH --job-name=piarr
#SBATCH --output=piarrJ.out      # can use %j and %A to separate
#SBATCH --partition=htc         # -p
#SBATCH --cpus-per-task=1       # do not need parallelism
#SBATCH --time=2:00
#SBATCH --mem=1G                # memory per node
#SBATCH --array=1-100%20

export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK:-1}

module load compilers/gcc-7.3.1
make piomp &>/dev/null

mtrials=${1:-1}
numtrials=$mtrials
seed=$((($RANDOM*$SLURM_ARRAY_TASK_ID))

#echo " /usr/bin/time -p ./piomp -t $numtrials"
./piomp -a -t $numtrials -s $seed | tee -a log
```